

# **Design and Analysis of Algorithms**

## **Part 3**

### **Data structures as a tool for algorithm design: heaps, heapsort, and priority queues**


Giulio Chiribella

Hilary Term 2019

# Heaps [CLRS 6.1]

A **heap** is a data structure that organizes data in an *essentially complete* rooted tree,  
i.e. a rooted tree that is **completely filled on all levels except possibly on the lowest, which is filled from the left up to a point.**

**Example:** binary heap, storing numbers (keys) at the nodes of the tree




The *height* of a tree is the longest simple path from the root to a leaf.

**Exercise.** Show that for a binary heap with  $n$  nodes, the height is  $\lfloor \lg n \rfloor$ .


# Implementation with arrays

A heap can be implemented by an array without any explicit pointers. In particular, a *binary* heap can be implemented by an array  $A$  as follows:

- Root of the binary tree is  $A[1]$
- Left child of  $A[i]$  is  $A[2i]$ .
- Right child of  $A[i]$  is  $A[2i + 1]$ .
- Hence, for  $i > 0$ , the parent of node  $i$  is the node  $\text{Parent}(i) = \lfloor i/2 \rfloor$ .



# Example



The heap is stored as the following array:

$$A = \begin{bmatrix} 6 & 4 & 22 & 8 & 11 & 9 & 3 & 32 & 7 & 1 \end{bmatrix}$$

**Exercise:** show that in a binary heap of  $n$  nodes there are  $\lfloor \frac{n+1}{2} \rfloor$  leaves, stored in the array elements with indices from  $\lceil \frac{n+1}{2} \rceil$  to  $n$ .

# Max-heaps

---

A *max-heap* is a heap that satisfies the

**Max-Heap Property:** The key of a node (except the root) is less than or equal to the key of its parent.


In the array implementation, the Max-Heap Property Reads:

*For all  $1 < i \leq A.heap\text{-}size$ :  $A[i] \leq A[\lfloor i/2 \rfloor]$ .*

## Remarks:

- The maximum element of a max-heap is at the root.
- In the following we will focus on *binary max-heaps*.  
Generally, a max-heap may be  $k$ -ary.
- One could also define *min-heaps*, where the key of each node (except the root) is larger than or equal to the key of its parent.

# Example



This is a max-heap. It can be stored in the array

$$A = \boxed{16 \mid 10 \mid 14 \mid 8 \mid 9 \mid 4 \mid 3 \mid 2 \mid 7 \mid 1}$$

Note that the array  $A$  is *not sorted*:

it does *not* satisfy the property  $A[i] \leq A[i - 1]$  for every  $i > 1$ .

However,  $A$  satisfies the max-heap property

$A[i] \leq A[\lfloor i/2 \rfloor]$  for every  $i > 1$ .

# Building a max-heap

---

Given an array  $A$ , there is a procedure to turn  $A$  into a max-heap:

**MAKE-MAX-HEAP**( $A$ )

Takes an array  $A$  of  $n$  integers and rearranges it into a max-heap of size  $n$ .

In turn, **MAKE-MAX-HEAP** is based on the following procedure:

**MAX-HEAPIFY**( $A, i$ )


*Assuming that the left and right subtrees of node  $i$  are max-heaps,*

**MAX-HEAPIFY** transforms the subtree rooted at the node  $i$  to a max-heap.

# MAX-HEAPIFY [CLRS 6.2]

**Idea:** compare the key at node  $i$  with the keys of its children, and rearrange them in order to satisfy the max-heap property.

**Example:**






# MAX-HEAPIFY [CLRS 6.2]

**Idea:** compare the key at node  $i$  with the keys of its children, and rearrange them in order to satisfy the max-heap property.




**Example:**



# MAX-HEAPIFY [CLRS 6.2]

**Idea:** compare the key at node  $i$  with the keys of its children, and rearrange them in order to satisfy the max-heap property.





**Example:**



# MAX-HEAPIFY [CLRS 6.2]

**Idea:** compare the key at node  $i$  with the keys of its children, and rearrange them in order to satisfy the max-heap property.

**Example:**



# MAX-HEAPIFY in pseudocode

---

MAX-HEAPIFY( $A, i$ )

**Input:** Assume left and right subtrees of  $i$  are max-heaps.

**Output:** Subtree rooted at  $i$  is a max-heap.

```
1   $l = 2i$                 //  $A[l]$  is the left-child of  $A[i]$ 
2   $r = 2i + 1$             //  $A[r]$  is the right-child of  $A[i]$ 
3  if  $l \leq n$  and  $A[l] > A[i]$  // Lines 4-8: Determine
4       $largest = l$         // largest among  $A[i]$ ,  $A[l]$  and  $A[r]$ .
5  else  $largest = i$ 
6  if  $r \leq n$  and  $A[r] > A[largest]$ 
7       $largest = r$ 
8  if  $largest \neq i$ 
9      exchange  $A[i]$  with  $A[largest]$ 
10     MAX-HEAPIFY( $A, largest$ )
```

# Running time of MAX-HEAPIFY

## MAX-HEAPIFY a subtree of size $n$ at node $i$

- $\Theta(1)$  to find the largest among  $A[i]$ ,  $A[2i]$  and  $A[2i + 1]$ .
- The subtree rooted at a child of node  $i$  has size upper bounded by  $2n/3$  (**Exercise.** Prove this fact.

Proof idea: the worst case is when last row of tree is exactly half full).

- Thus  $T(n) \leq T(2n/3) + \Theta(1)$ .
- By the Master Theorem, we have

$$T(n) = O(n^0 \log n) = O(\log n).$$

## Alternative reasoning:

Define the *height* of a node to be the number of edges on the longest simple downward path from the node to a leaf.


On a node of height  $h$ , MAX-HEAPIFY runs for  $O(h)$  time at most.

The height of the root of a heap of size  $n$  is  $\lfloor \lg n \rfloor$ , so  $T(n) = O(\log n)$ .

# MAKE-MAX-HEAP [CLRS 6.3]

**Idea:** starting from the last *non-leave* node, apply MAX-HEAPIFY to the subtree based at that node. Repeat the same procedure for all the previous nodes.


## Example



# MAKE-MAX-HEAP [CLRS 6.3]

**Idea:** starting from the last *non-leave* node,  
apply MAX-HEAPIFY to the subtree based at that node.  
Repeat the same procedure for all the previous nodes.


## Example



# MAKE-MAX-HEAP [CLRS 6.3]

**Idea:** starting from the last *non-leave* node, apply MAX-HEAPIFY to the subtree based at that node. Repeat the same procedure for all the previous nodes.

## Example







# MAKE-MAX-HEAP [CLRS 6.3]

**Idea:** starting from the last *non-leave* node,  
apply MAX-HEAPIFY to the subtree based at that node.  
Repeat the same procedure for all the previous nodes.


## Example



# MAKE-MAX-HEAP (example continued)



## MAKE-MAX-HEAP (example continued)



Note that the procedure works because at every step the left and right subtrees are max-heaps.

# Pseudocode

---

Recall that the leaves are the nodes indexed by  $\lceil \frac{n+1}{2} \rceil, \lceil \frac{n+1}{2} \rceil + 1, \dots, n$ .

MAKE-MAX-HEAP( $A$ )

**Input:** An (unsorted) integer array  $A$  of length  $n$ .

**Output:** A heap of size  $n$ .

```
1   $A.heap\text{-}size = A.length$ 
2  for  $i = \lceil \frac{n+1}{2} \rceil - 1$  downto 1
3      MAX-HEAPIFY( $A, i$ )
```

# Correctness

---

**Loop invariant:** *Each node  $i + 1, i + 2, \dots, n$  is the root of a max-heap.*

## Initialization

Each node  $\lceil \frac{n+1}{2} \rceil, \lceil \frac{n+1}{2} \rceil + 1, \dots, n$  is a leaf, which is the root of a trivial max-heap. Since  $i = \lceil \frac{n+1}{2} \rceil - 1$  before the first iteration, the invariant is initially true.

## Maintenance

Suppose  $i = i_0 \geq 1$  and assume each node  $i_0 + 1, i_0 + 2, \dots, n$  is the root of a max-heap. Executing  $\text{MAX-HEAPIFY}(A, i)$  causes  $i_0$  to be the root of a new max-heap. Hence each node  $i_0, i_0 + 1, \dots, n$  is now the root of a max-heap, meaning that the loop invariant holds after  $i$  has been decremented from  $i_0$  to  $i_0 - 1$ .

## Termination

When  $i = 0$  (i.e. after the counter becomes less than 1) the loop terminates. By the loop invariant, each node, in particular node 1, is the root of a max-heap.

# Running time analysis

**Simple (but loose) bound:**  $O(n \log n)$ .

We have  $O(n)$  calls to MAX-HEAPIFY, each taking  $O(\log n)$  time.

**Tighter analysis:**  $O(n)$ .

MAX-HEAPIFY takes linear time in the height of the node it runs on, and “most nodes have small heights”.

*Fact.* The number of nodes of height  $h$  is upper bounded by  $n/2^h$ , and the cost of MAX-HEAPIFY on a node of height  $h$  is  $\leq ch$ , for some  $c > 0$ .

Hence, the cost of MAKE-MAX-HEAP is

$$T(n) \leq \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{n}{2^h} ch \leq cn \left( \sum_{h=0}^{\infty} \frac{h}{2^h} \right) = 2cn,$$

*Note.* Differentiating  $\sum_{k=0}^{\infty} x^k = \frac{1}{1-x}$  and multiplying by  $x$ , we get  $\sum_{k=0}^{\infty} kx^k = \frac{x}{(1-x)^2}$ , for  $|x| < 1$ .

# Applications of heaps

---

- Sorting: *heapsort*, an *in-place* sorting algorithm with worst-case complexity  $O(n \log n)$ .
- Efficient implementation of *priority queues*:  
Max-heap  $\rightarrow$  max-priority queue.  
Min-heap  $\rightarrow$  min-priority queue.  
Max-priority queues can be used to schedule jobs on a shared computer.  
Min-priority queues can be used to simulate events in time.

**Remark.** Actual implementations often have a *handle* in each heap element that allows access to an object in the application, and objects in the application often have a handle (likely an array index) to access the heap element.

# Heapsort [CLRS 6.4]

---

A sorting algorithm based on the heap data structure.

**Idea.** Given an input array,

- Build a max-heap using MAKE-MAX-HEAP.
- Starting from the root (maximum element), place the maximum element into the correct place in the array by swapping it with the element in the last position in the array.
- “Discard” this last node – decrement the heap size, and call MAX-HEAPIFY on the smaller structure with the possibly incorrectly-placed root.
- Repeat this discarding process until only one node (the minimum) remains, and is therefore in the correct place in the array.


**Features:**

- $O(n \log n)$  worst case – like merge sort.
- Sorts *in place* – like insertion sort.





# Example: heapsort

---






# Example: heapsort

---







# Example: heapsort






---









# Example: heapsort









# Example: heapsort



# Example: heapsort




# Example: heapsort



4      7







4      7

4      7



4      7



# Example: heapsort



4      7

4      7

4      7



4      7

4      7



# The algorithm

---

HEAPSORT( $A$ )

```
1  MAKE-MAX-HEAP( $A$ )
2  for  $i = A.heap\text{-}size$  downto 2
3      exchange  $A[1]$  with  $A[i]$ 
4       $A.heap\text{-}size = A.heap\text{-}size - 1$ 
5      MAX-HEAPIFY( $A, 1$ )
```

## Running time

- MAKE-MAX-HEAP takes  $O(n)$
- The for-loop is executed  $n - 1$  times.
- Exchange operation takes  $O(1)$ .
- MAX-HEAPIFY takes  $O(\log n)$ .

Total time:  $O(n \log n)$ .

# Priority queues [CLRS 6.5]

---

A Priority queue is an *abstract data structure* for maintaining a set of elements, each with an associated value called a *key*.

Max-priority queues give priority to the elements with larger keys, min-priority queues give priority to the elements with smaller keys.

## Operations supported by a max-priority queue:

1.  $\text{INSERT}(S, x, k)$  inserts element  $x$  with key  $k$  into set  $S$ .
2.  $\text{MAXIMUM}(S)$  returns the element of  $S$  with the largest key.
3.  $\text{EXTRACT-MAX}(S)$  removes and returns the element of  $S$  with the largest key.
4.  $\text{INCREASE-KEY}(S, x, k)$  increases value of  $x$ 's key to  $k$ .  
Requires  $k$  to be at least as large as  $x$ 's current key value.

# Priority queues [CLRS 6.5]

A Priority queue is an *abstract data structure* for maintaining a set of elements, each with an associated value called a *key*.

Max-priority queues give priority to the elements with larger keys, min-priority queues give priority to the elements with smaller keys.

## Operations supported by a max-priority queue:

1.  $\text{INSERT}(S, x, k)$  inserts element  $x$  with key  $k$  into set  $S$ .
2.  $\text{MAXIMUM}(S)$  returns the element of  $S$  with the largest key.
3.  $\text{EXTRACT-MAX}(S)$  removes and returns the element of  $S$  with the largest key.
4.  $\text{INCREASE-KEY}(S, x, k)$  increases value of  $x$ 's key to  $k$ .  
Requires  $k$  to be at least as large as  $x$ 's current key value.

**Operations supported by a min-priority queue** supports  $\text{INSERT}(S, x)$ ,  $\text{MINIMUM}(S)$ ,  $\text{EXTRACT-MIN}(S)$  and  $\text{DECREASE-KEY}(S, x, k)$ .

# Implementation by unordered-sequence

---

Store the elements  $e$  and their keys  $k$  (as pairs  $(e, k)$ ) in an unordered sequence, implemented as an array or a *doubly-linked list*.

- Implement  $\text{INSERT}(S, e, k)$  by inserting  $(e, k)$  at the end of the sequence; takes  $O(1)$  time.
- Implement  $\text{EXTRACT-MAX}(S)$  by inspecting all elements of the sequence and removing the maximum; takes  $\Theta(n)$  time.

*We can do better with a heap implementation!*

# Implementation by heap

---

- A heap offers a good compromise between insertion and extraction. Both operations take  $O(\log n)$  time.
- For simplicity, in the following, we identify the element with its key.

# Implementation by heap

---

- A heap offers a good compromise between insertion and extraction. Both operations take  $O(\log n)$  time.
- For simplicity, in the following, we identify the element with its key.

## Finding the maximum

HEAP-MAXIMUM( $A$ )

**return**  $A[1]$

*Time:*  $\Theta(1)$

# Extracting maximum

---

- Check that the heap is non-empty.
- Make a copy of the maximum element (root).
- Make the last node in the tree the new root.
- HEAPIFY the array, but *less the last node*.
- Return the copy of the maximum.

HEAP-EXTRACT-MAX( $A$ )

```
1  if  $A.heap-size < 1$ 
2      error “heap underflow”
3   $max = A[1]$ 
4   $A[1] = A[A.heap-size]$ 
5   $A.heap-size = A.heap-size - 1$ 
6  MAX-HEAPIFY( $A, 1$ )
7  return  $max$ 
```

*Time:*  $O(\log n)$ , where  $n$  is the size of the heap.

# Increasing key value

---

Given set  $S$ , entry  $i$ , and new key value  $key$ :

1. Check that  $key$  is greater than or equal to  $i$ 's current value.
2. Update  $i$ 's key value to  $key$ .
3. Traverse the tree upward comparing  $i$  to its parent and swapping keys if necessary, until  $i$ 's key is smaller than its parent's key.

HEAP-INCREASE-KEY( $A, i, key$ )

```
1  if  $key < A[i]$ 
2      error “new key is smaller than current key”
3   $A[i] = key$ 
4  while  $i > 1$  and  $A[\text{Parent}(i)] < A[i]$ 
5      exchange  $A[i]$  with  $A[\text{Parent}(i)]$ 
6       $i = \text{Parent}(i)$ 
```

*Time.*  $O(\log n)$



# Insertion

---

Given a key  $k$  to insert into the heap:

- Insert a new node in the very last position in the tree with key  $-\infty$ .
- Increase the  $-\infty$  key to  $k$  using HEAP-INCREASE-KEY

HEAP-INSERT( $A, key$ )

- 1  $A.heap-size = A.heap-size + 1$
- 2  $A[A.heap-size] = -\infty$
- 3 HEAP-INCREASE-KEY( $A, A.heap-size, key$ )

*Time.*  $O(\log n)$