

# **Design and Analysis of Algorithms**

## **Part 6**

### **Paths in Graphs**

Giulio Chiribella

Hilary Term 2019

# How to find the optimal route?

---

In many applications, like GoogleMaps, it is important to find the shortest route from one location to another.

More generally, going from one location to another could have a *cost* and it is important to find the route that minimizes the cost.

## Graph formulation:

- locations  $\rightarrow$  vertices
- connections between two locations  $\rightarrow$  edges
- cost of going from one location to another  $\rightarrow$  weight on the edge
- route  $\rightarrow$  path
- cost of a route  $\rightarrow$  sum of the weights of the edges on the path.

# Shortest paths

Consider a directed graph  $G = (V, E)$  with **weight function** (or *length*)  $w : E \longrightarrow \mathbb{R}_{\geq 0}$ .

□ The **weight** (or *length*) **of a path**  $p = \langle v_0, v_1, \dots, v_k \rangle$  is

$$w(p) := \sum_{i=1}^k w(v_{i-1}, v_i)$$

I.e.  $w(p)$  is the sum of edge weights on path  $p$ .

□ The **shortest-path weight** (or *length*) from vertex  $u$  to vertex  $v$  is

$$\delta(u, v) := \begin{cases} \min\{ w(p) : p \text{ is a path from } u \text{ to } v \} & \text{if } \exists \text{ path from } u \text{ to } v \\ \infty & \text{otherwise} \end{cases}$$

□ A **shortest path** from  $u$  to  $v$  is a path  $p$  such that  $w(p) = \delta(u, v)$ .

# Breadth-first search [CLRS 22.2]

---

Consider the simplest case: *all weights equal to 1*.

$\delta(u, v)$  = minimum number of edges on a path from  $u$  to  $v$ , if such a path exists; otherwise,  $\delta(u, v) = \infty$ .

Given a source vertex  $s$ , *breadth-first search (BFS)* finds all the vertices reachable from  $s$  and, for each reachable vertex  $v$  it finds the shortest-path length  $\delta(s, v)$  and a shortest path from  $s$  to  $v$ .

**Input:** A graph  $G = (V, E)$ , either directed or undirected,  
and a source vertex  $s \in V$ .

**Output:** For each  $v \in V$ , an integer  $d[v]$  and a back pointer  $\pi[v]$ , such that

1.  $d[v] = \delta(s, v)$
2.  $\pi[v] = u$  is the predecessor of  $v$  on a shortest path from  $s$  to  $v$ .  
If there is no path from  $s$  to  $v$ , then  $\pi[v] = \text{NIL}$ .

# Intuitive idea

---

**Idea:** Send a wave out from the source  $s$ .

- The wave first hits all vertices 1 edge away from  $s$ .
- From there, the wave will then hit all vertices 2 edges away from  $s$ .
- When the wave reaches  $v$ , we know the distance from  $s$  to  $v$ .  
The shortest path is constructed by keeping track of the vertices reached by the wave on the way to  $v$ .

**Data structure implementation:** the wave can be implemented by a *first-in-first-out (FIFO) queue*  $Q$ .

A vertex  $v$  is put in the queue when it is reached by the wave, and it is removed by the queue after all its neighbours are reached.

# FIFO queues [CLRS 10.1]

---

A **FIFO queue** is an *abstract data structure* with three basic operations:

1. ENQUEUE( $Q, x$ ): Inserts  $x$  at the *end* of the queue  $Q$ .
2. DEQUEUE( $Q$ ): Returns and removes the item at the *head* of the queue  $Q$ .
3. ISEMPTY( $Q$ ): Returns whether or not the queue is empty.

In the pseudo code we will use the notation  $Q = \emptyset$  and  $Q \neq \emptyset$ .

## Implementation:

Linked list with an extra “END” pointer to the tail of the list.

The head of the queue corresponds to the head of the linked list.

This makes all operations  $O(1)$ .

# The BFS algorithm [CLRS 22.2]

BFS( $V, E, s$ )

**Input:** A directed or undirected graph  $(V, E)$ , and a source  $s \in V$ .

**Output:** For each  $v \in V$ ,  $d[v]$  and  $\pi[v]$

```
1   $d[s] = 0; \pi[s] = nil$            // Source can be reached with path of length 0
2  for each  $u \in V - \{s\}$          // Mark all other nodes
3       $d[u] = \infty$                // as not reached yet
4       $\pi[u] = nil$ 
5   $Q = \emptyset$                    // Initialise  $Q$ 
6  ENQUEUE( $Q, s$ )                 // to contain only source  $s$ 
7  while  $Q \neq \emptyset$ 
8       $u = \text{DEQUEUE}(Q)$          // Node  $u$  is finished
9      for each  $v \in \text{Adj}[u]$ 
10         if  $d[v] = \infty$         // If  $v$  is reached for the first time.
11              $d[v] = d[u] + 1$     // record shortest distance,
12              $\pi[v] = u$           // create backpointer,
13             ENQUEUE( $Q, v$ )      // and join queue of unfinished nodes
```

# Running time

---


BFS takes  $O(|V| + |E|)$  in total.

- $O(|V|)$  because of the initialization step and of the enqueue/dequeue operations (each vertex is enqueued/dequeued at most once)
- $O(|E|)$  because we examine the edge  $(u, v)$  only when  $u$  is dequeued. Hence every edge is examined at most once (or at most twice if the graph is undirected and we represent edges as unordered pairs  $\{u, v\}$ ).



# Example

Consider the following graph with source  $s = A$ .



First, the wave hits vertices  $\{ B, C, F \}$ .

Second, it hits  $\{ E, D, G \}$ .

Third, it hits  $\{ H \}$ .

At the fourth step, the wave does not discover any new nodes.


# The BFS tree

The **BSF tree**\* consists of the **vertices reachable from  $s$**  and has an **edge between  $u$  and  $v$  iff  $u = \pi[v]$** .

Explicitly, the BSF tree is the graph  $G_\pi := (V_\pi, E_\pi)$  with  $V_\pi := \{u \in V : \pi[u] \neq \text{NIL}\} \cup \{s\}$  and  $E_\pi := \{(\pi[u], u) : u \in V_\pi\}$ .

\* note the usual abuse of notation: in math, trees are *undirected graphs*.

## Example (cont'd)



**Note:**  $d[v]$  is the level of  $v$  in the BFS tree.

**Note:** the BFS tree depends on the order in which adjacent nodes are explored, for example, if  $E$  came before  $D$ , the edge  $(E, H)$  would be in the tree instead of  $(D, H)$ .

## Lower bound on $d[v]$

**Lemma 1** If  $d[v] < \infty$ , there exists a path of length  $d[v]$  from  $s$  to  $v$ .

**Proof.** By induction on  $d[v]$ .

*Base case.* If  $d[v] = 0$ , then  $v$  must be the source  $s$ . Hence, Lemma 1 trivially holds.

*Induction step.* Suppose that, for every  $v$  with  $d[v] \leq d_0$ , there exists a path of length  $d[v]$  from  $s$  to  $v$ . Suppose that  $v$  is such that  $d[v] = d_0 + 1$ . Then, let  $u = \pi[v]$  be the predecessor of  $v$ . By construction  $d[v] = d[u] + 1$ , and therefore  $d[u] = d_0$ .

Now, the induction hypothesis guarantees that there exists a path of length  $d[u]$  from  $s$  to  $u$ . Adding the edge  $(u, v)$  to this path, we obtain a path of length  $d[u] + 1 = d[v]$  from  $s$  to  $v$ .  $\square$

**Corollary 1 (lower bound on  $d$ ).**  $d[v] \geq \delta(s, v)$  for every vertex  $v$ .

## Properties of the queue [Lemma 22.3, p. 599 of CLRS]

**Lemma 2** Let  $Q = \langle v_1, \dots, v_r \rangle$  be the queue at a given step of BFS. Then,  $d[v_r] \leq d[v_1] + 1$  and  $d[v_i] \leq d[v_j]$  for every  $i$  and  $j \in \{1, \dots, r\}$ .

**Proof.** By induction on the total number  $n$  of ENQUEUE/DEQUEUE operations. *Base case.* For  $n = 1$ , only vertex  $s$  is been enqueued, and Lemma 2 trivially holds.

*Induction step.* Suppose that Lemma 2 holds for  $n \leq n_0$  queue operations. If a DEQUEUE operation is performed, Lemma 2 still holds. If a ENQUEUE operation is performed, let  $v$  be the vertex that is being enqueued, and let  $u = \pi[v]$  be its predecessor. Note that  $d[u] \leq d[v_1]$ , since there are two possibilities: either (1)  $\pi[v_1] = u$ , in which case  $d[u] = d[v_1] - 1$ , or (2)  $\pi[v_1] \neq u$ , in which case  $u$  and  $v_1$  must have been both in the queue at earlier time, and induction implies  $d[u] \leq d[v_1]$ . Now, ENQUEUE sets  $v_{r+1} := v$  and  $d[v_{r+1}] := d[u] + 1$ . Hence,  $d[v_{r+1}] = d[u] + 1 \leq d[v_1] + 1$ . It remains to show  $d[v_r] \leq d[v_{r+1}]$ . Again, there are two possibilities: (1)  $\pi[v_r] = u$ , in which case  $d[v_r] = d[u] + 1 = d[v_{r+1}]$ , and (2)  $\pi[v_r] \neq u$ , in which case  $u$  and  $v_r$  must have been both in  $Q$  at an earlier time, and induction implies  $d[v_r] \leq d[u] + 1 = d[v_{r+1}]$ .

## Upper bound on $d[v]$

**Corollary 2.** If  $u$  is enqueued before  $v$ , then  $d[u] \leq d[v]$ .

**Proof.** Immediate from Lemma 2.

**Lemma 3 (upper bound on  $d$ ).** If there exists a path of length  $l$  from  $s$  to  $v$ , then  $d[v] \leq l$ .

**Proof.** By induction on  $l$ . *Base cases.* For  $l = 0$ ,  $v = s$  and  $d[s] = 0$ . For  $l = 1$ , the path is an edge  $(s, v)$ , and therefore  $d[v] = 1$ .

*Induction step.* Suppose that Lemma 3 holds for paths of length  $l \leq l_0$ , and suppose that  $\langle s, v_1, \dots, v_{l_0}, v \rangle$  is a path of length  $l_0 + 1$  from  $s$  to  $v$ . By the induction hypothesis,  $d[v_{l_0}] \leq l_0$ . Now, there are three possibilities:

(1)  $v$  is enqueued before  $v_{l_0}$  is enqueued. Then, Corollary 2 implies  $d[v] \leq d[v_{l_0}] \leq l_0$ .

(2)  $v$  is enqueued before  $v_{l_0}$  is dequeued. Then, Lemma 2 implies  $d[v] \leq d[v_{l_0}] + 1 \leq l_0 + 1$ .

(3)  $v$  is enqueued after  $v_{l_0}$  is dequeued. Then,  $\pi[v] = v_{l_0}$  and  $d[v] = d[v_{l_0}] + 1 = l_0 + 1$ . □

# Correctness of BFS

---

**Theorem 1.**  $d[v] = \delta(s, v)$  for every vertex  $v \in V$ .

**Proof.** Suppose that  $\delta(s, v)$  is finite. Then, there exists a path of length  $\delta(s, v)$  from  $s$  to  $v$ . Hence, Lemma 3 yields  $d[v] \leq \delta(s, v)$ . Since we also have  $d[v] \geq \delta(s, v)$  (Corollary 1), we conclude  $d[v] = \delta(s, v)$ .

Suppose that  $\delta(s, v)$  is infinite. Then, Corollary 1 implies  $d[v] \geq \delta(s, v) = \infty$ . □

## Correctness of BFS (cont'd)

**Theorem 2.** If  $0 < d[v] < \infty$ , then  $\pi[v]$  is the predecessor of  $v$  on a shortest path from  $s$  to  $v$ .

**Proof.** By induction on  $d[v]$ . *Base case.* For  $d[v] = 1$ ,  $(s, v)$  is an edge in  $E$  and, by construction,  $\pi[v] = s$ .

*Induction step.* Suppose that Theorem 2 holds for every vertex  $v$  such that  $d[v] \leq l_0$ . For a vertex  $w$  with  $d[w] = l_0 + 1$ , define  $v_{l_0} := \pi[w]$ . By construction,  $d[v_{l_0}] = d[w] - 1 = (l_0 + 1) - 1 = l_0$ .

Hence, Lemma 1 implies that there exists a path  $\langle s, v_1, \dots, v_{l_0} \rangle$ . Adding the edge  $(v_{l_0}, w)$ , we obtain the path  $p = \langle s, v_1, \dots, v_{l_0}, w \rangle$ .

The length of the path is  $l_0 + 1 = d[w]$ . Hence, Theorem 1 implies that  $p$  is a shortest path from  $s$  to  $w$ . The predecessor of  $w$  in  $p$  is  $v_{l_0} = \pi[w]$ .  $\square$

# DFS vs BFS [DPV 4.2]

---

- DFS goes “as far as possible” (depth), while BSF discovers all the vertices at a certain distance before moving to further vertices.  
In terms of data structures, this is reflected in the fact that DFS uses (implicitly) a last-in-first-out queue (i.e. a stack), while BSF uses a first-in-first-out queue.
- Unlike DFS, BFS *may not reach all vertices*:  
BFS discovers only the vertices that are reachable from the source  $s$ .
- Both algorithms have linear running time  $O(|V| + |E|)$ .



# Dijkstra's Algorithm [CLRS 24.3]

---

- Dijkstra's algorithm solves the *Single-Source Shortest Path Problem* for *non-negative weights*.
- It is essentially a weighted version of breath-first search.
- The algorithm uses a *min-priority queue*  $Q$  (instead of FIFO queue), with keys given by the *shortest-path weight estimates*  $d[v]$ .
- At termination,
  - $d[v]$  is equal to the distance from  $s$  to  $v$
  - the algorithm provides a back-pointer  $\pi[v]$ , such that  $\pi[v]$  is the predecessor of  $v$  on a shortest path from  $s$  to  $v$ , if such path exists, or  $\pi[v] = \text{NIL}$  otherwise.

# Pseudocode



DIJKSTRA( $V, E, w, s$ )

**Input:** A directed or undirected graph  $(V, E)$ ,  $s \in V$ ,  $w : E \rightarrow \mathbb{R}_{\geq 0}$ .

**Output:** For each  $v \in V$ ,  $d[v]$  and  $\pi[v]$

```
1  for each  $v \in V$ 
2       $d[v] = \infty$ 
3       $\pi[v] = nil$ 
4   $d[s] = 0$ 
5   $Q = \text{MAKE-QUEUE}(V)$  with  $d[v]$  as keys
6  while  $Q \neq \emptyset$ 
7       $u = \text{EXTRACT-MIN}(Q)$            // Shortest route to  $u$  known
8      for each vertex  $v \in \text{Adj}[u]$ 
9          if  $d[u] + w(u, v) < d[v]$     // Shorter route to  $v$  via  $u$  discovered
10              $d[v] = d[u] + w(u, v)$ 
11              $\pi[v] = u$ 
12              $\text{DECREASE-KEY}(Q, v, d[v])$ 
```

# Example



$n$ -th iteration	Init	1	2	3	4	5
EXTRACT-MIN( $Q$ )		$A$	$C$	$B$	$E$	$D$
(source) $d[A]$	0	0	0	0	0	0
$d[B]$	$\infty$	4	3	3	3	3
$d[C]$	$\infty$	2	2	2	2	2
$d[D]$	$\infty$	$\infty$	6	5	5	5
$d[E]$	$\infty$	$\infty$	7	4	4	4
$Q$	$\langle A, B, C, D, E \rangle$	$\langle C, B, D, E \rangle$	$\langle B, D, E \rangle$	$\langle E, D \rangle$	$\langle D \rangle$	$\langle \rangle$

# Correctness

---

## Loop invariants for the **while** loop:

- I1.** *For all  $v \in V$ , we have  $d[v] \geq \delta(s, v)$ .*
- I2.** *For all  $v \in S := V \setminus Q$ , we have  $d[v] = \delta(s, v)$ .*

## Initialisation:

- I1.** Just before the start of the **while** loop,  $d[s] = 0 = \delta(s, s)$   
and  $d[v] = \infty$  for every other vertex  $v \in V, v \neq s$ .  
Hence I1 holds.
- I2.** Before the start of the **while** loop  $S = \emptyset$ , so I2 is vacuously true.

**Termination:** At the end,  $S = V$ .

Hence, I2 implies that  $d[v] = \delta(s, v)$  for all  $v \in V$ , as required.

## Maintenance of I1: $d[v] \geq \delta(s, v) \quad \forall v \in V$

Suppose that I1 holds before an iteration of the **while** loop.

We have to show that I1 still holds after the iteration.

For a generic vertex  $v \in V$ , there are two possibilities:  
either  $d[v]$  did not change after the iteration, or  $d[v]$  changed.

1. If  $d[v]$  did not change, then  $d[v] \geq \delta(s, v)$  by I1 before the iteration.
2. If  $d[v]$  changed, let  $u$  be the vertex selected by EXTRACT-MIN in this iteration of the loop. Then, we have the inequality

$$\begin{aligned} d[v] &= d[u] + w(u, v) \\ &\geq \delta(s, u) + w(u, v) \quad \text{by I1 before the iteration} \\ &\geq \delta(s, v) \end{aligned},$$

where the last inequality holds because  $\delta(s, u) + w(u, v)$  is the length of a path from  $s$  to  $v$ .  $\square$

**Terminology:**  $\delta(s, v) \leq \delta(s, u) + w(u, v)$  is called the *triangle inequality*.

## Maintenance of I2: $d[v] = \delta(s, v)$ , $\forall v \in S$

---

Suppose that I2 holds before an iteration of the **while** loop.


We have to show that I2 still holds after the iteration.

For a generic vertex  $v \in S$  there are three possibilities:

1.  $v$  was in  $S$  before the iteration of the **while** loop.  
Then,  $d[v] = \delta(s, v)$  by I2 before the iteration.
2.  $v$  has been added to  $S$  in this iteration *and* there is no path from  $s$  to  $v$ .  
Then, the  $d$ -value is  $d[v] \geq \delta(s, v) = \infty$ , namely  $d[v] = \infty$ .
3.  $v$  has been added to  $S$  in this iteration *and* there is a path from  $s$  to  $v$ .  
Then, there is also a *shortest* path  $s \xrightarrow{p} v$ .  
Just before  $v$  is added to  $S$ , the path  $p$  connects a vertex in  $S$  (the source  $s$ ) with a vertex outside  $S$  (vertex  $v$ ).  
Let  $y$  be the first vertex along  $p$  that is outside  $S$ ,  
and let  $x$  be the predecessor of  $y$  (by definition,  $x$  is in  $S$ ).

...picture on the next slide.

## Maintenance of I2: $d[v] = \delta(s, v)$ , $\forall v \in S$ (cont'd)



**Optimal substructure:** since  $s \xrightarrow{p_1} x \rightarrow y \xrightarrow{p_2} v$  is a shortest path to  $v$ ,  
 $s \xrightarrow{p_1} x \rightarrow y$  is a shortest path to  $y$ .

Since  $x$  was in  $S$  before the iteration of the **while** loop,  
we have  $d[x] = \delta(s, x)$  by I2 before the iteration.

Moreover, when  $x$  was extracted from  $Q$ ,  
the  $d$ -value of  $y$  must have been updated by lines 10-11 of the code.

...continue on the next slide.

# The convergence property

## Lemma 1 (Convergence property).

*Suppose  $s \xrightarrow{p_1} x \rightarrow y$  is a shortest path.*

*If  $d[x] = \delta(s, x)$  before lines 10-11 are executed on vertex  $y$ ,  
then  $d[y] = \delta(s, y)$  after lines 10-11 are executed on vertex  $y$ .*

After executing lines 10 and 11, we have


$$\begin{aligned} d[y] &\leq d[x] + w(x, y) && \text{by lines 10-11} \\ &= \delta(s, x) + w(x, y) && \text{by hypothesis} \\ &= \delta(s, y) && \text{because } s \xrightarrow{p_1} u \rightarrow v \text{ is a shortest path} \end{aligned}$$

By I1, we have  $d[y] \geq \delta(s, y)$ .

Hence  $d[y] = \delta(s, y)$ . □



## Maintenance of I2: $d[v] = \delta(s, v)$ , $\forall v \in S$ (cont'd)



The Convergence Property implies  $d[y] = \delta(s, y)$  before  $v$  is extracted from  $Q$ .

Since  $y$  and  $v$  were both in  $Q$  when  $v$  was chosen by EXTRACT-MIN, we have  $d[v] \leq d[y] = \delta(s, y) \leq \delta(s, v)$ .

In conclusion, we obtained  $d[v] \leq \delta(s, v)$ .

Since  $d[v] \geq \delta(s, v)$  by I1, we obtained  $d[v] = \delta(s, v)$ .

This proves the validity of I2 after the iteration of the **while** loop. □

# Constructing shortest paths

---

**Theorem.** For  $v \neq s$  and  $d[v] < \infty$ ,  $\pi[v]$  is the predecessor of  $v$  on a shortest path from  $s$  to  $v$ .

**Proof.** Let  $u := \pi[v]$ , and let  $s \xrightarrow{p} u$  be a shortest path from  $s$  to  $u$ . At termination, we have

$$\begin{aligned}\delta(s, v) &= d[v] && \text{by I2} \\ &= d[u] + w(u, v) && \text{because } u = \pi[v] \\ &= \delta(s, u) + w(u, v) && \text{by I2}\end{aligned}\tag{1}$$

Hence  $s \xrightarrow{p} u \rightarrow v$  is a shortest path from  $s$  to  $v$ . □

# Running time

---

The running time of Dijkstra's algorithm depends on the implementation of the min-priority queue. For a min-heap implementation, one has

1. MAKE-QUEUE:  $O(|V|)$
2. EXTRACT-MIN:  $O(\log |V|)$
3. DECREASE-KEY:  $O(\log |V|)$

MAKE-QUEUE is executed once.

EXTRACT-MIN is executed  $|V|$  times.

Since each vertex  $v \in V$  is added to  $S$  exactly once, each edge in  $Adj[v]$  is examined in the **for** loop exactly once.

Thus there are a total of  $|E|$  iterations of the **for** loop, and hence a total of at most  $|E|$  DECREASE-KEY operations.

**Total running time:**

$$\begin{aligned} & O(|V|) + O(|V| \log |V|) + O(|E|) + O(|E| \log |V|) \\ = & O((|V| + |E|) \log |V|) \end{aligned}$$

# A variation: Unique Shortest Path

---

**Question:** is shortest path to a given vertex is unique or not?

We will now see an algorithm that answers this question.

**Input:** A directed or undirected graph  $G = (V, E)$   
with non-negative weights  $w : E \rightarrow \mathbb{R}_{\geq}$ ,  
and a source vertex  $s \in V$ .

**Output:** A Boolean array  $usp[v]$  such that, for each vertex  $v$ ,  
 $usp[v] = \text{TRUE}$  iff the shortest path from  $s$  to  $v$  is unique.

Note:

- $usp[s] = \text{TRUE}$
- if  $\delta(s, v) = \infty$  we conventionally assign the value  $usp[v] = \text{TRUE}$

# Pseudocode for the Unique Shortest Path Problem

---

The problem can be solved by modification of Dijkstra's algorithm.

The array  $usp[-]$  is initialized to TRUE.

The **while** loop is modified as follows:


```
1  while  $Q \neq \emptyset$ 
2       $u = \text{EXTRACT-MIN}(Q)$ 
3      for each vertex  $v \in \text{Adj}[u]$ 
4          if  $d[u] + w(u, v) < d[v]$ 
5               $d[v] = d[u] + w(u, v)$ 
6               $\pi[v] = u$ 
7               $usp[v] = usp[u]$ 
8               $\text{DECREASE-KEY}(Q, v, d[v])$ 
9          else if  $d[v] = d[u] + w(u, v)$ 
10              $usp[v] = \text{FALSE}$ 
```

# Shortest paths in graphs with negative weights

So far, we assumed that the weight function  $w : E \rightarrow \mathbb{R}$  was non-negative, namely that  $w(u, v) \geq 0$  for every edge  $(u, v) \in E$ .

*How about graphs where some of the weights are negative?*

## Problem: negative-weight cycles



$\langle B, C, B \rangle$  is a negative-weight cycle:  $w(B, C) + w(C, B) = -2$ .  
There is *no shortest path* from source  $A$  in this graph!

# General graphs: Bellman-Ford Algorithm [CLRS 24.1]

BELLMAN-FORD( $V, E, w, s$ )

**Input:** A directed or undirected graph  $(V, E)$ ,  $s \in V$ ,  $w : E \rightarrow \mathbb{R}$ .

**Output:** FALSE, if there exist a negative-weight cycle reachable from  $s$ ,  
otherwise TRUE, and for each  $v \in V$ ,  $d[v]$  and  $\pi[v]$ .

```
1  for each  $v \in V$ 
2       $d[v] = \infty; \pi[v] = nil$ 
3   $d[s] = 0$ 
4  for  $i = 1$  to  $|V| - 1$            // correctly computes the distances
5      for each edge  $(u, v) \in E$     // if there is no negative-weight
6          if  $d[u] + w(u, v) < d[v]$   // cycle reachable from  $s$ 
7               $d[v] = d[u] + w(u, v); \pi[v] = u$ 
8  for each edge  $(u, v) \in E$         // checks for negative-weight cycles
9      if  $d[u] + w(u, v) < d[v]$       // reachable from  $s$ 
10         return FALSE
11     else return TRUE
```

# Running time

---

- Initialisation takes  $\Theta(|V|)$ .
- Each of the  $|V| - 1$  iterations of lines 5–7 takes  $\Theta(|E|)$
- the **for** loop of lines 8-11 takes  $O(|E|)$ .

In total, the running time is  $\Theta(|V||E|)$ .



# Correctness: the case of no negative-weight cycles

If no negative-weight cycle is reachable from  $s$ ,  
then the shortest paths are well-defined.

**Invariant of the for loop 4-7.** Before the  $i$ -th iteration of the loop,

**I1**  $d[v] \geq \delta(s, v), \forall v \in V$

**I2** *for every shortest path  $\langle v_0, v_1, \dots, v_k \rangle$  from  $v_0 = s$  to  $v_k$ ,*  
 $d[v_j] = \delta(s, v_j), \forall j < i$

**Initialisation.** Before the first iteration,  $d[s] = 0 = \delta(s, s)$ , and  $d[v] = \infty \geq \delta(s, v)$  for every  $v \in V, v \neq s$ . Hence, I1 holds. Moreover,  $d[v_0] = d[s] = 0 = \delta(s, s)$ . Hence, I2 holds.

**Termination.** Termination occurs for  $i = |V|$ .

Then, I2 implies  $d[v_j] = \delta(s, v_j)$  for all  $j < |V|$ . Since there are  $|V|$  vertices in the graph, one has  $k < |V|$ , and therefore  $d[v_j] = \delta(s, v_j)$  for all  $j < k$ . Hence, if there is a path from  $s$  to  $v$ ,  $d[v] = \delta(s, v)$ . If there is no path,  $\delta(s, v) = \infty$  and  $d[v] \geq \delta(s, v)$  by I1, whence  $d[v] = \delta(s, v)$ .

## Maintenance of I1: $d[v] \geq \delta(s, v), \forall v \in V$

---

**Fact:** I1 is maintained also by the *internal* **for** loop (lines 5-7).

Indeed, consider the iteration of the internal **for** loop on the edge  $(u, v)$

□ Only the  $d$ -value of  $v$  can change in this iteration.

For all the other  $d$ -values, I1 still holds.

□ If  $v$  did not change, then I1 still holds.

□ If  $d[v]$  changed, then  $d[v] = d[u] + w(u, v)$ .

Hence, one has

$$\begin{aligned} d[v] &= d[u] + w(u, v) \\ &\geq \delta(s, u) + w(u, v) \quad \text{by I1 before the iteration} \\ &\geq \delta(s, v). \end{aligned}$$

Since I1 is maintained by the internal **for** loop 5-7, it is also maintained by the external **for** loop 4-7.

## Maintenance of I2: $d[v_j] = \delta(s, v_j), \forall j < i$

Assume that I2 holds before the  $i$ -th iteration of the **for** loop 4-7.

**Fact:** For every  $v \in V$ , the  $d$ -value  $d[v]$  does not increase after each iteration of the internal **for** loop 5-7 (and therefore it does not increase after each iteration of the external **for** loop 4-7)

For every  $j < i$ , the new  $d$ -value  $d'[v_j]$  is no larger than the old  $d$ -value  $d[v_j] = \delta(s, v_j)$ . By I1, this implies  $d'[v_j] = \delta(s, v_j)$ .

For  $j = i$ , consider the iteration of the internal **for** loop on the edge  $(v_{i-1}, v_i)$ . After this iteration, one has

$$\begin{aligned} d[v_i] &\leq d[v_{i-1}] + w(v_{i-1}, v_i) \\ &= \delta(s, v_{i-1}) + w(v_{i-1}, v_i) && \text{because I2 holds for all } j < i \\ &= \delta(s, v_i) && \text{because } \langle s, v_1, \dots, v_k \rangle \text{ is a shortest path,} \\ &&& \text{whence } \langle s, v_1, \dots, v_i \rangle \text{ is a shortest path} \end{aligned}$$

Since the  $d$ -values are non-increasing, this bound continues to hold after other iterations of the internal **for** loop. Hence,  $d[v_i] = \delta(s, v_i)$ .

# Constructing shortest paths

---

**Theorem.** Suppose that no negative-weight cycle is reachable from  $s$ . Then, for  $v \neq s$  and  $d[v] < \infty$ ,  $\pi[v]$  is the predecessor of  $v$  on a shortest path from  $s$  to  $v$ .

**Proof.** Same proof as in Dijkstra's algorithm.

# Correctness: the case of negative-weight cycles

**Theorem.** If there is a negative-weight cycle reachable from  $s$ , then the **for** loop 5-7 will return FALSE.

**Proof.** Suppose that there is a negative weight-cycle reachable from  $s$ , say  $c = (v_1, v_2, \dots, v_n, v_{n+1})$  with  $v_{n+1} = v_1$ , and  $w(c) := \sum_{i=1}^n w(v_i, v_{i+1}) < 0$ .

**Fact:** we have the inequality

$$\begin{aligned} \sum_{i=1}^n \left[ d[v_i] + w(v_i, v_{i+1}) \right] &= \left( \sum_{i=1}^n d[v_i] \right) + w(c) \\ &< \sum_{i=1}^n d[v_i] \\ &= \sum_{i=1}^n d[v_{i+1}]. \end{aligned}$$

Since the l.h.s. is strictly smaller than the r.h.s., there must exist a vertex  $v_i$  such that  $d[v_i] + w(v_i, v_{i+1}) < d[v_{i+1}]$ . Hence, the **for** loop of lines 8-11 will return FALSE upon inspecting the edge  $(v_i, v_{i+1})$ . □

# Shortest paths in DAGs [CLRS 24.2]

**Idea:** Find a topological sort of  $V$ . The shortest path from  $s$  to  $v$  can only contain intermediate vertices that come *before*  $v$  in the topological sort.

**Input:** A DAG  $G = (V, E)$  with weight  $w : E \rightarrow \mathbb{R}$   
and a source vertex  $s$

**Output:** For each  $v \in V$ , length  $d[v]$  of shortest-path from  $s$  to  $v$

```
1  TOPOLOGICAL-SORT( $V, E$ )
2  for each  $v \in V$ 
3       $d[v] = \infty$ 
4       $\pi[v] = nil$ 
5   $d[s] = 0$ 
6  for each  $u \in V$  in topological order
7      for each  $v \in Adj[u]$ 
8          if  $d[v] > d[u] + w(u, v)$ 
9               $d[v] = d[u] + w(u, v)$ 
10              $\pi[v] = u$ 
```

# Running time

---

1.  $\text{TOPOLOGICAL-SORT}(V, E)$  takes  $O(|V| + |E|)$  time
2. initialization (lines 2-5) takes  $\Theta(|V|)$  time
3. line 6 is run one time on each vertex, resulting into a  $O(|V|)$  time
4. the **for** loop of line 7 runs  $O(|E|)$  time

**Total running time:**  $O(|V| + |E|)$

# Correctness

## Invariant of the **for** loop of lines 6-10.

**I1**  $d[v] \geq \delta(s, v)$  for all  $v \in V$

**I2** Let  $(v_1, v_2, \dots, v_n)$  be the topological sort.  
Before the **for** loop is executed on vertex  $v_k$ ,  
one has  $d[v_i] = \delta(s, v_i)$  for every  $i \leq k$ .

## Initialisation:

□  $d[s] = \delta(s, s)$ . For  $v \neq s$ ,  $d[v] = \infty \geq \delta(s, v)$ . Hence, I1 holds.

□ if  $v_1 = s$ ,  $d[v_1] = 0 = \delta(s, s)$  before the loop starts.

Hence, I2 holds for  $v_1 = s$

□ if  $v_1 \neq s$ , then there cannot be a path from  $s$  to  $v_1$ .

Hence,  $\delta(s, v_1) = \infty = d[v_1]$  and I2 holds for  $v_1 \neq s$ .

**Termination.** At termination, I2 yields  $d[v] = \delta(s, v)$  for every vertex  $v \in V$ .

**Maintenance of I1:** proof as usual (cf. proof in Bellman-Ford algorithm).



## Maintenance of I2: $d[v_i] = \delta(s, v_i)$ for every $i \leq k$ .

---

When the **for** loop is executed on vertex  $v_k$ , only vertices in  $Adj[v_k]$  are affected. That is, only vertices  $v_j$  with  $j > k$  are affected (because of topological sort). Hence, we only need to consider vertex  $v_{k+1}$ .

Let  $v_j$  be the predecessor of  $v_{k+1}$  on a shortest path from  $s$  to  $v_{k+1}$ .

By construction,  $j < k + 1$ , and therefore  $d[v_j] = \delta(s, v_j)$  by I2 before iteration.

After the **for** loop 6-10 was executed on  $v_j$ , one had

$$\begin{aligned} d[v_{k+1}] &\leq d[v_j] + w(v_j, v_{k+1}) \\ &= \delta(s, v_j) + w(v_j, v_{k+1}) \quad \text{by the invariant before iteration} \\ &= \delta(s, v_{k+1}). \end{aligned}$$

Since further iterations of the **for** loop can not increase the  $d$ -value, and since I1 holds, we conclude that  $d[v_{k+1}] = \delta(s, v_{k+1})$  after the **for** loop is executed on  $v_k$ .

# Constructing shortest paths

---

**Theorem.** For  $v \neq s$  and  $d[v] < \infty$ ,  $\pi[v]$  is the predecessor of  $v$  on a shortest path from  $s$  to  $v$ .

**Proof.** Same proof as in Dijkstra's algorithm.

# All-pairs shortest paths [CLRS 25]

## All-pairs shortest paths

**Input:** A directed graph  $(V, E)$  with general weights  $w : E \rightarrow \mathbb{R}$ ,  
and **with no negative-weight cycles**.

**Task:** For each pair of vertices  $u$  and  $v$ , find shortest path from  $u$  to  $v$ .

## Using Single-source algorithms

- Using Bellman-Ford algorithm for each vertex, leads to a  $O(|V|^2|E|)$  algorithm, which is  $O(|V|^3)$  for a sparse graph ( $|E| = O(|V|)$ ) and  $O(|V|^4)$  for a dense graph ( $|E| = O(|V|^2)$ ).
- If there are no negative-weight edges, using Dijkstra's algorithm leads to a  $O((|V| + |E|) |V| \log |V|)$  algorithm.

## Using dynamic programming

- $O(|V|^3)$  with the Floyd-Warshall algorithm (next slides).

# The Floyd-Warshall algorithm [CLRS 25.2]

---

Suppose the vertex-set is  $\{0, \dots, n-1\}$  and let

$$d[i, j; k] = \begin{cases} \text{length of shortest path from } i \text{ to } j, \text{ all of whose} \\ \text{intermediate nodes are in the set } \{0, \dots, k-1\} \end{cases}$$

Initially

$$d[i, j; 0] = \begin{cases} w(i, j) & \text{if } (i, j) \in E \\ \infty & \text{otherwise} \end{cases}$$

# Floyd-Warshall algorithm, cont'd

Suppose that we have calculated  $d[i, j; k]$  for every  $i$  and  $j$ , and we want to compute  $d[i, j; k + 1]$ .

**Optimal substructure:** if  $s \xrightarrow{p_1} u \xrightarrow{p_2} v$  is a shortest path, then  $s \xrightarrow{p_1} u$  and  $u \xrightarrow{p_2} v$  are shortest paths.

**Fact:** since there are no negative cycles (by assumption), a shortest path from  $i$  to  $j$  that uses intermediate vertices in  $\{0, \dots, k\}$  goes through  $k$  at most once.

Hence, we have

$$d[i, j; k + 1] = \min \left\{ d[i, j; k], d[i, k; k] + d[k, j; k] \right\}$$

# Pseudocode

---

FLOYD-WARSHALL( $V, E, w$ )

```
1  for  $i = 0$  to  $|V| - 1$ 
2      for  $j = 0$  to  $|V| - 1$ 
3           $d[i, j; 0] = \infty$ 
4  for each edge  $(i, j) \in E$ 
5       $d[i, j; 0] = w(i, j)$ 
6  for  $k = 0$  to  $|V| - 1$ 
7      for  $i = 0$  to  $|V| - 1$ 
8          for  $j = 0$  to  $|V| - 1$ 
9               $d[i, j; k + 1] = \min \left\{ d[i, j; k], d[i, k; k] + d[k, j; k] \right\}$ 
```

**Running time** is  $O(|V|^3)$ .

# Summary of Shortest Paths Algorithms

- *Breadth-First-Search* solves the **single-source** shortest path problem for directed or undirected graphs with **unit weight**, i.e. where  $w(e) = 1$  for all  $e \in E$  in  $O(|V| + |E|)$ .
- *Dijkstra's Algorithm* solves the **single-source** shortest-path problem for directed or undirected graphs with **non-negative weights** in  $O((|V| + |E|) \log |V|)$ .
- The *Bellman-Ford Algorithm* solves the **single-source** shortest-path problem for directed or undirected graphs with **arbitrary weights** in  $O(|V||E|)$ . In addition, it checks **checks for negative-weight cycles**.
- *Shortest Paths in DAGs* solves the **single-source** shortest path problem with **arbitrary weights** in the case of a **directed acyclic graph** in  $\Theta(|V| + |E|)$ .
- The *Floyd-Warshall Algorithm* solves the **all-pairs** shortest path problem for directed or undirected graphs with **no negative-weight cycles** in  $O(|V|^3)$ .