

Imperative Programming 3

Lecture 10: Implementation

Peter Jeavons

Trinity Term 2019



Implementing OOP

Aim: understand how the main features of object-oriented languages are implemented and how programs are organised in **memory**

We'll first look at these features:

- **Identity**: each instance of a class has a distinct identity
- **Polymorphism**: classes with same *interface* can be used interchangeably
- **Inheritance**: class may be defined as extension of another

Identity

Object identity is simply the **address** of the chunk of memory allocated to the object

- Not to be confused with the hash code which is printed
- Good enough for C++ (user-driven memory management)
- In Scala/Java the **garbage collector** dynamically *moves objects around* to remove unused objects
 - Keeps the free memory contiguous
 - Objects *change addresses* dynamically
 - What if another object is using one which has moved?
- Hence Scala/Java make it impossible to get address


Objects in memory

```
class Car {  
  private var reg: Int  
  private var miles: Int
```

```
  ...
```

```
  def drive(dist: Int) {  
    this.miles = this.miles + dist  
  }  
}
```

```
val c = new Car
```



- Each object has a (shared) pointer to a **class descriptor**
- Instance variables allocated at fixed offsets from base
- Each method has implicit parameter this
- The miles field has the address this+8

Polymorphism

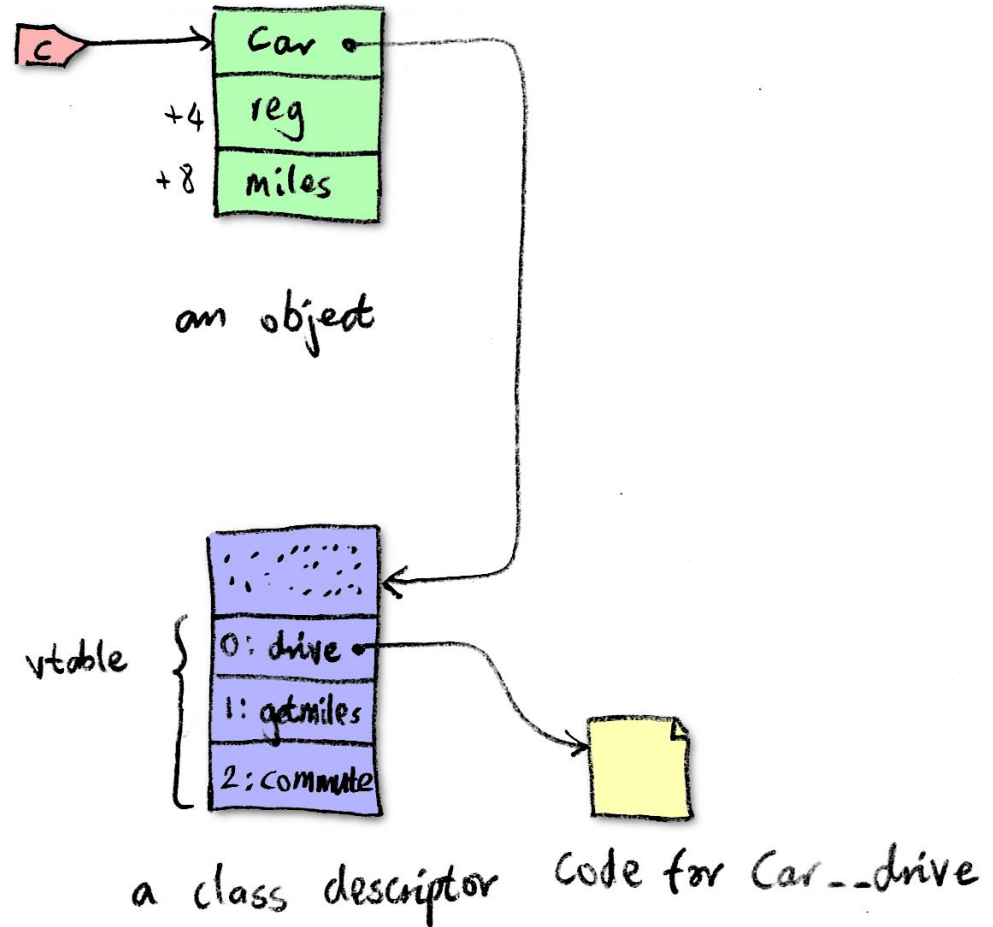
- Recall that which method is called on an object is determined by the *dynamic type* not the *static type*

```
var v: Vehicle  
var c: Car  
v=c  
v.drive(100)
```

- Here the v is a parent *trait* (or class)
- Compiler verifies static type (i.e. that Vehicle has drive)
- Runtime uses the dynamic type (drive method in Car)

```
v.drive(100)
```

Method invocation



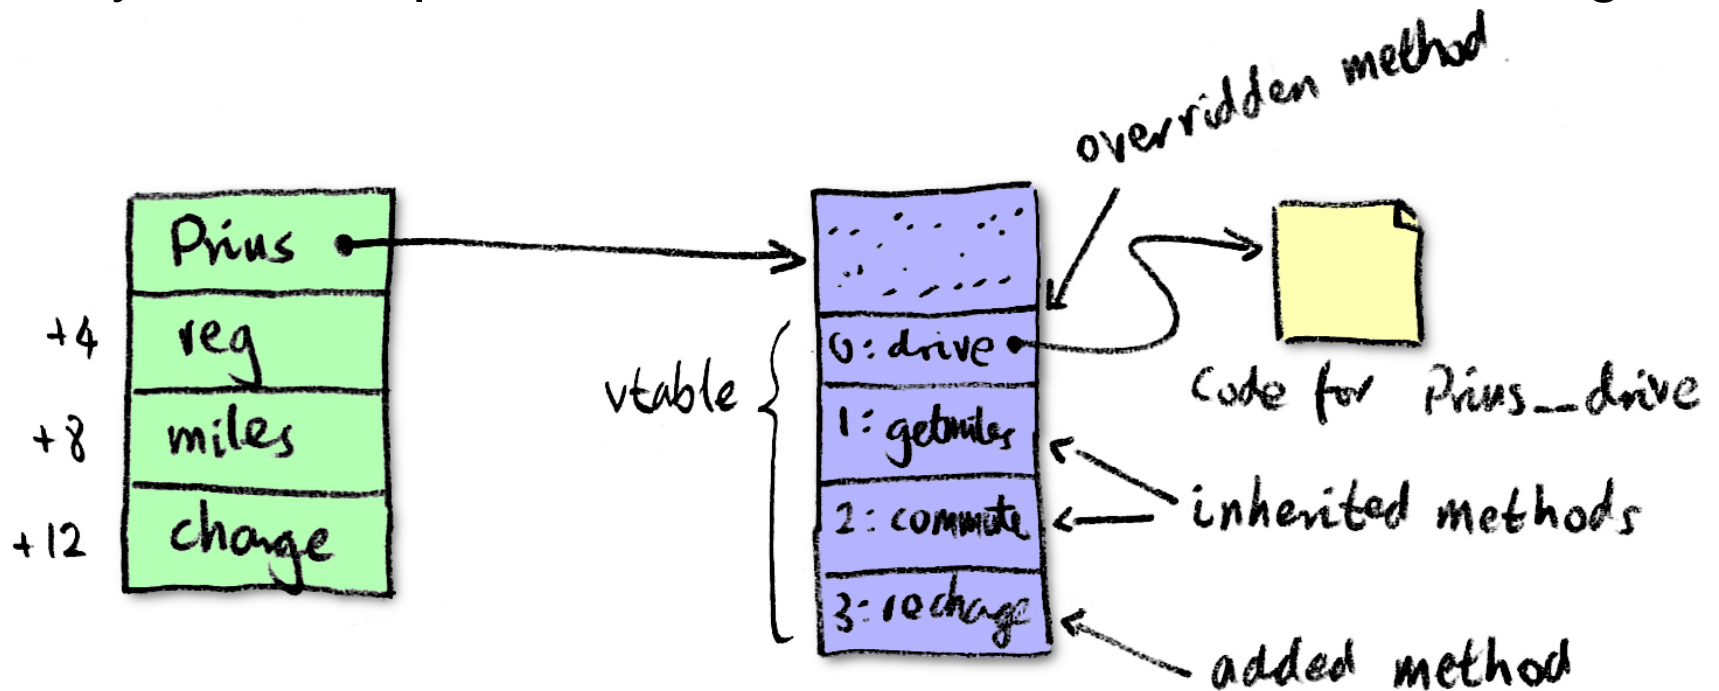
```
v.drive(100)
```

is shorthand for

```
v.class.vtable[0](v, 100)
```

Inheritance

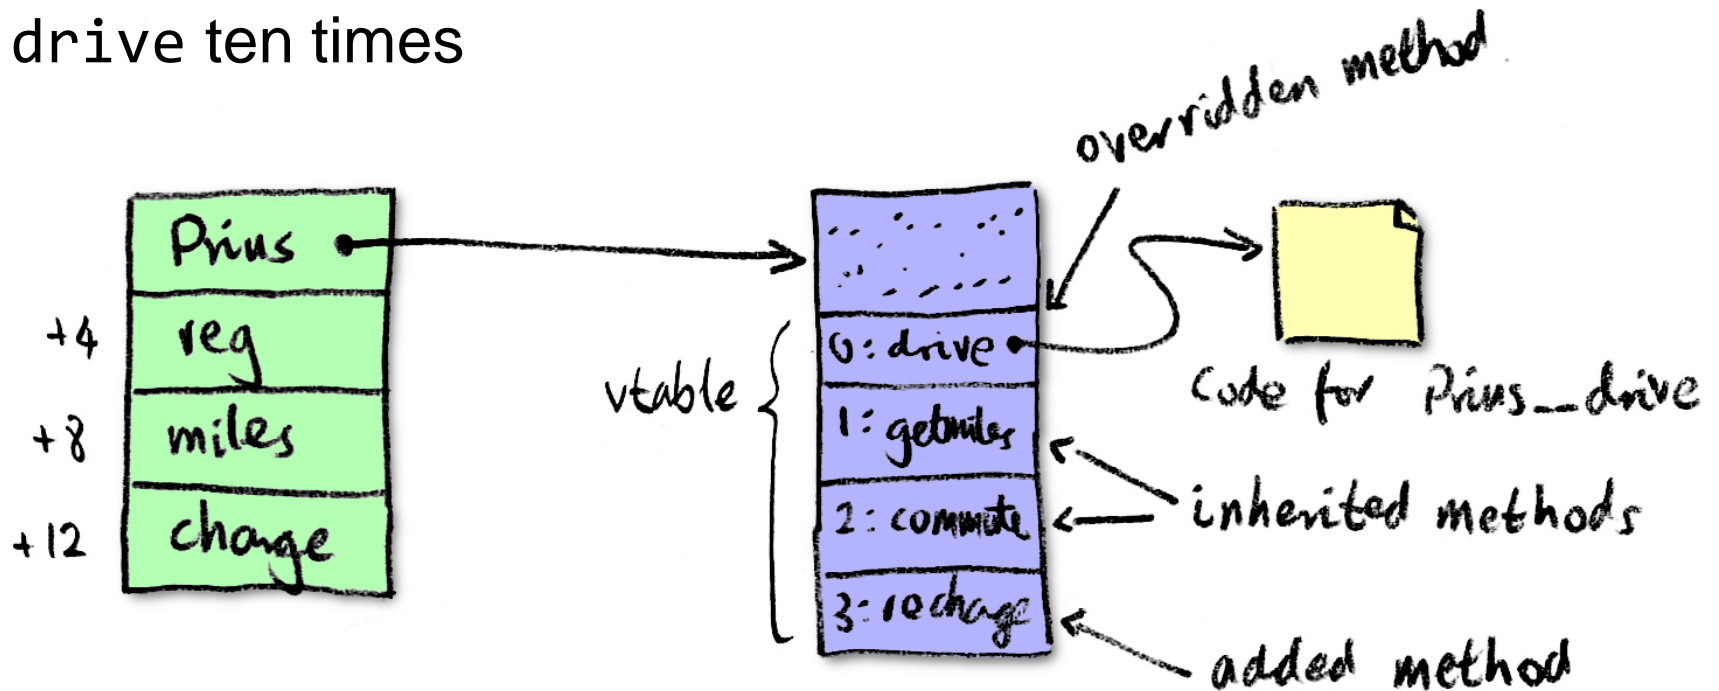
- In a subclass: extend superclass' object record and vtable
- Instance variables are inherited (and extended)
- Methods are inherited, extended or overridden
- Dynamic dispatch relies on consistent method indexing



Late binding of methods

```
// In class Car
def commute() {
  for (i <- 1 to 10) this.drive(50)
}
```

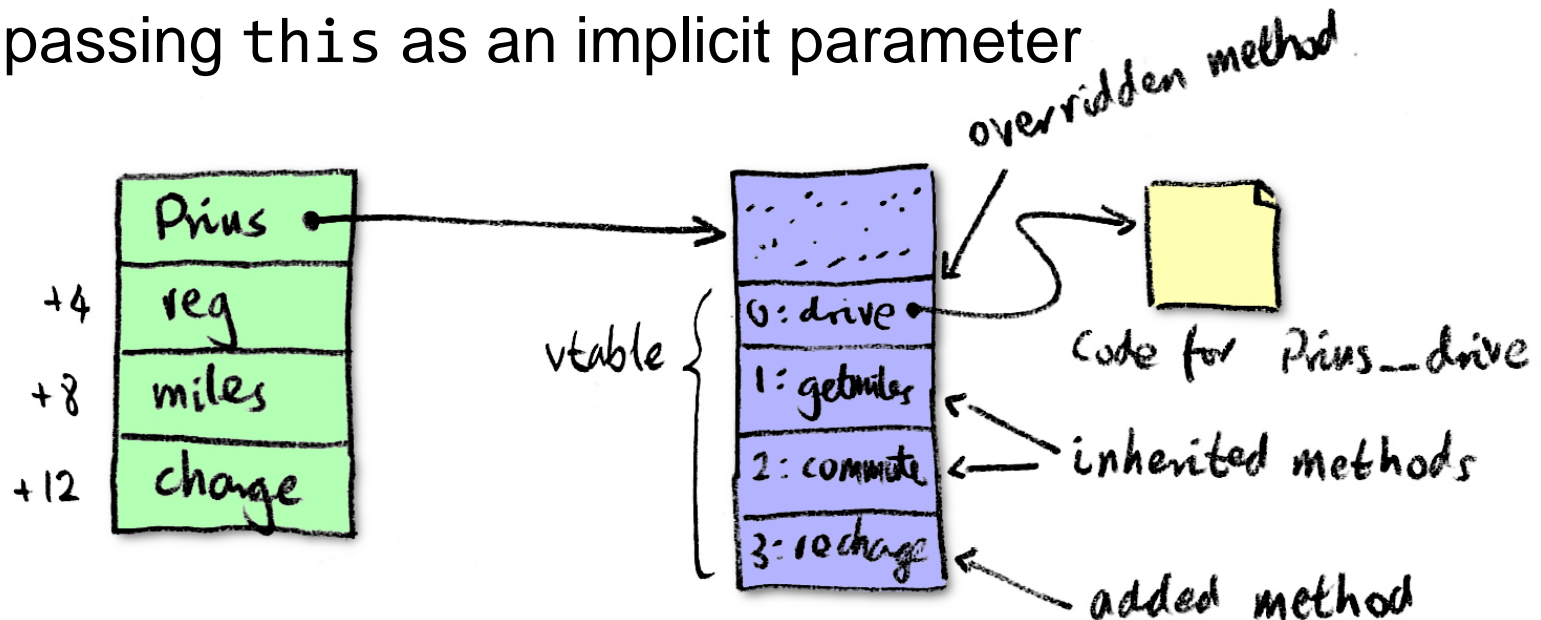
Consequence: calling commute on a Prius object will invoke commute in the Car superclass and call the Prius version of drive ten times



Super calls

```
// In class Prius
override def drive(dist: Int) {
  super.drive(dist)
  this.charge = this.charge - dist/10
}
```

- This drive method behaves as before & reduces charge
- 'Super call' unambiguously refers to Car version of drive passing this as an implicit parameter



Some problems

- The *fragile binary interface problem*
 - If we change Car class then its record or vtable layout may change
 - Changes in offsets mean that all subclasses of Car should be recompiled
 - In Scala the Java Virtual Machine (JVM) delays layout until the program is loaded which adds an overhead
- The *virtual lookup problem*
 - Each virtual call requires a least one extra dereference which adds an overhead at runtime
 - Can be mitigated by an optimising compiler inlining some virtual calls (thus avoiding a vtable lookup)
 - Can also be mitigated by the programmer e.g. avoid overridden method calls inside tight loops and instead place the loop inside the method
- Main message: advanced language features have a runtime cost...
 - ...but possibly the cost is **not so high** as some claim

Memory management

Memory management is the process of **allocating** new objects and **removing** unused objects to make space for those new object allocations.

- **Memory-related bugs** are among the most prevalent and difficult to catch of all software bugs
 - especially, in programs written in an unsafe language (C/C++)
 - but can also arise in Java or Scala programs

Storing Objects in Main Memory

- Objects must be stored somewhere
 - in **main memory**
 - storing objects in secondary storage is a separate problem
- Memory **allocation strategy**:
 - When/how is memory allocated/released?
 - Who takes care of these actions?
 - programmer vs. the “language”
- Three common types of memory allocation:
 - static
 - stack-based
 - heap-based

Static Allocation

- Memory is reserved at **compile time**
 - constants (e.g., arrays) are stored in the executable
 - memory is allocated when program starts
 - memory is released when program terminates
- Main benefit: efficiency
- Drawbacks:
 - program can use only a fixed amount of memory
 - inflexible for representing complex structures
 - prevent recursion
- Available in most languages
 - Fortran, C, C++, Ada

Static Memory Allocation in Java

- Class fields can be made static
 - they exist only once per class (AKA **singletons**)
 - i.e., different objects share the same field
- Example: want to count how many Region objects have been allocated

```
public class Region {  
    protected static int regionsNo;  
    static {  
        regionsNo = 0;  
    }  
  
    public Region(...) {  
        regionsNo++;  
        ...  
    }  
}
```

Only one regionsNo field exists for all instances of Region.

Static initializer: code that is executed when a class is loaded to initialize the class.

Static Memory Allocation in Scala

- **Companion objects** - Scala's equivalent of static
- The companion object of a class
 - has the same name as the class
 - is defined in the same file as the class

```
object Region {  
  private var regionsNo = 0  
  def count = regionsNo  
}
```

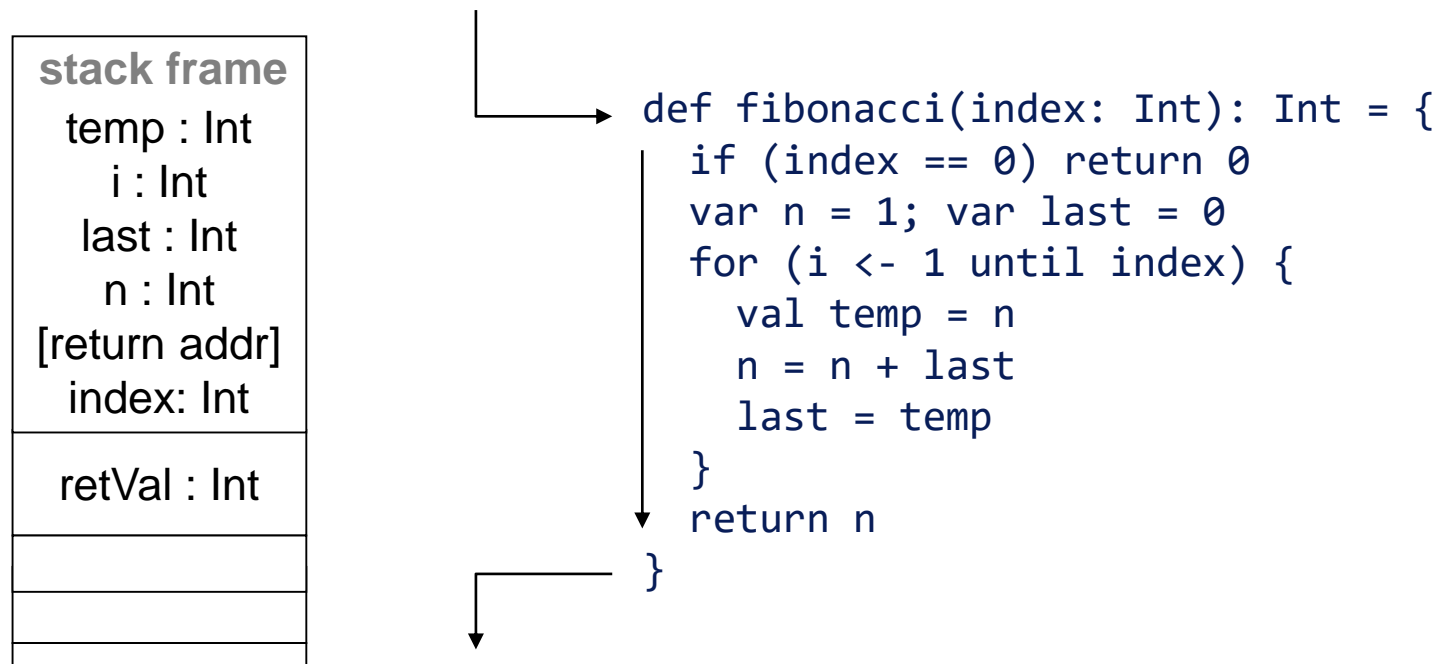
```
class Region(...) {  
  Region.regionsNo += 1  
}
```

← The object and class can access each other's private fields and methods.

← Access to the fields and methods in the companion object must be qualified with the name of the object.

Stack-based Allocation

- Each thread of execution has a **stack of frames**
- A frame is allocated when a method is called
 - Contains **return address**, **parameters**, **local variables**, and **return value**
- The frame is popped off the stack when the method ends



Stack-based Allocation


- Some OO languages (e.g., C++) can store objects on the stack
 - objects are constructed when a method starts
 - objects are destroyed when a method ends
 - using references to objects on stack causes problems
 - references become invalid when a method terminates
- Java/Scala cannot store objects on the stack
 - avoids the problem of references to objects on stack
 - stack frames in Java/Scala...
 - ... store only local variables
 - ... cannot be manipulated directly by the programmer

Heap-based Allocation

- Main way of allocating memory in Java/Scala
- **Heap**: global memory pool
 - maximum size determined by -Xmx JVM option
- Memory is allocated **as needed** using new
- But when to release memory?
 - upon termination
 - **must** release memory as it becomes unused
 - programs such as servers will otherwise run out of memory
- Two approaches to memory reclamation:
 - explicit (i.e., manual) reclamation
 - automatic garbage collection

Manual Memory Reclamation

- In some languages, unused objects must be removed from memory **explicitly**
 - e.g., in C++, one can write `delete b2.author`



- If we `delete b2.author`, we **must** update `b1.author` as well
- If we forget, then `b1.author` points to an **invalid object**
 - AKA dangling reference
 - reading from `b1.author` returns an arbitrary value
 - writing into `b1.author` results in a run-time error


the most common source of bugs in practice!

Automatic Garbage Collection

- Alternative: release memory **automatically** when it becomes unused
 - typically done by the **virtual machine** or the **runtime library**
 - AKA **(automatic) garbage collection**
- Benefits:
 - no problems with dangling references
 - no need to worry about when objects become unused
 - less error-prone
- Main drawback: efficiency
 - overhead in determining which objects are unused
 - unsuitable for systems with critical performance
- Fundamental question: what is **unused memory**?

Reference Counting

- Main idea:
 - associate a counter `o.cnt` with each object `o`
 - increase `o.cnt` when attaching `o` to a reference
 - decrease `o.cnt` when detaching `o` from a reference
 - if `o.cnt` is zero, then `o` is unused



Now `o1.cnt = 0`, so `o1` is unused and its memory can be reclaimed.

```
r1 = new Object()
```


- create `o1`
- assign `o1` to `r1`
- increase `o1.cnt`

```
r1 = new Object()
```

- create `o2`
- decrease `o1.cnt`
- assign `o2` to `r1`
- increase `o2.cnt`

Limits of Reference Counting

```
class A {  
    var ref: A = null  
}
```




```
r1 = new A()  
r1.ref = new A()  
r1.ref.ref = r1  
r1 = null
```

- Both o1 and o2 are unused,...
- ...but their counters are 1
 - the memory of o1 and o2 never get released

Reference counting cannot deal with **cyclic references**

Unreachable Objects

- An object is **unused** if it cannot be reached through references starting from the **active code**
 - i.e., starting from the references on the stack
 - to ensure an object is not reachable, set the reference to null
- **Mark-and-sweep** garbage collection:
 - stop the execution of all threads
 - phase 1: mark all objects that are reachable from the active code
 - phase 2: release all unmarked objects



Neither object is reachable from r1

⇒ Neither object is marked

⇒ Both objects can be reclaimed

When is Garbage Collection Run?

- JVM offers no guarantees!
 - can be source of performance problems:
 - the program needs to do something promptly, but the JVM is running garbage collection
- Garbage collection is **typically** run when...
 - ...the available heap memory is exhausted
 - ...the program is idle
- Can be requested manually
 - `System.gc()`

Memory Leaks

- **Memory leak**: unused memory is not reclaimed due to programmer's oversight
 - common problem in C++ and other languages with explicit memory mgmt
 - But can this happen in languages with automatic memory management (e.g., in Java or Scala)?
- **Example**:
 - creating two Region objects with the same (x, y, width, height) wastes memory
 - solution: **internalize** Region
 - create a global cache of Region objects
 - reuse the object with the same (x, y, width, height)
 - internalization is a common **memory usage optimization** technique
 - internalized objects are still value objects
 - Java strings provide a similar mechanism

Internalizing Region

```
object Region {  
  protected val regions = new mutable.HashMap[Region, Region]()  
}
```

```
class Region(...) {
```

```
  // ... must implement equals() and hashCode()
```

```
  def internalize(): Region =  
    Region.regions.get(this) match {  
      case Some(r) => r  
      case None => {  
        regions(this) = this  
        this  
      }  
    }  
}
```

← regions is common to all instances of Region

- Looks for an equal object in regions
- If not found, then registers itself there
- Returns the object from regions

Internalization and Memory Leaks

```
val r1 = new Region(10, 10, 40, 40).internalize()
val r2 = new Region(10, 10, 40, 40).internalize()
assert(r1 eq r2) // AnyRef.eq - reference equality
```

- Assertion holds!
 - internalization guarantees that each object exists only once


- But internalized objects are never freed!

```
r1 = null
r2 = null
```

- releasing explicit references does not help:
 - regions holds an active reference to the region (10, 10, 40, 40)
 - static objects are reachable from active code
- \Rightarrow With time, internalization uses ever more memory
 - memory leak!
- Solutions:
 - do nothing – acceptable if memory exhaustion is unlikely
 - use **weak references**

Weak References

- **Weak reference**: does not prevent an object from being garbage collected
- Example:



- o1 and o2 are not reachable through **strong references**
⇒ they are unused and can be freed

Internalization w/o Memory Leaks

```
object Region {  
  protected val regions =  
    new mutable.WeakHashMap[Region, WeakReference[Region]]()  
}
```

← java.lang.ref.WeakReference

```
class Region(...) {
```

```
  // ... must implement equals() and hashCode()
```

```
  def internalize(): Region =
```

```
    Region.regions.get(this) match {
```

```
      case Some(r) if r.get != null => r.get
```

```
      case _ => {
```

```
        regions(this) = new WeakReference[Region](this)
```

```
        this
```

```
      }
```

```
    }
```

← Checks if the object has not been garbage collected

Source of Memory Leaks in Java/Scala

- Memory leaks occur when objects are added to **long-lived** collections
 - static collections
 - collections that exist as long as the application
 - associated with windows and GUI elements
 - improper usage or bugs in third-party libraries
 - improper usage of the Observer pattern
 - common in GUI applications
- Memory leaks can occur in **native code**

Finalizers

- `finalize()`:
 - called before an object is garbage collected
 - provides an opportunity for cleanup
- Example: files must be closed after use

```
class FileInputStream extends InputStream {  
    ...  
    override def finalize(): Unit = {  
        if (...) {    // checks whether the file is open  
            close()  
        }  
    }  
}
```

- without the finalizer, a `FileInputStream` object could be garbage-collected without releasing the resources
 - results in **resource leak**!

Nondeterministic Nature of Finalizers

```
object A {  
  private var numFinalized = 0  
  
  def main(args: Array[String]) = {  
    for (i <- 0 until 1000000) {  
      val a = new A  
    }  
    System.out.println(A.numFinalized)  
  }  
}  
class A {  
  override def finalize(): Unit = {  
    A.numFinalized += 1  
  }  
}
```

Five runs on my machine:

571222

562831

522754

270059

554158

Releasing Resources


- **Do not** rely on finalizers to release resources!
 - finalizers **are not destructors** from C++!
 - no guarantee when or if finalizers are invoked
- Release resources explicitly as soon as possible
 - e.g., close a file as soon as you are done using it
 - add an explicit method for releasing resources
 - e.g., `close()`, `dispose()`
- Add finalizers whenever possible as an “insurance” against **careless programmers**

Releasing Resources and Exceptions

- Resources are often not released due to exceptions
 - solution: add a finally block

If an exception is thrown here, os is not closed!

```
val os: FileOutputStream = ...  
os.write(...)  
os.close()
```



close() is called even if an exception is thrown.

```
val os: FileOutputStream = ...  
try {  
    os.write(...)  
}  
finally {  
    os.close()  
}
```

Summary

- Memory is divided into **static**, **stack** and **heap**
- **Java/Scala** objects are stored on heap
- Manual memory reclamation is error-prone
- **Garbage Collection** automatizes memory reclamation
 - detects and disposes unreachable objects
 - but cannot prevent memory leaks
- **Weak references** suitable for caches
- Releasing resources as soon as not needed